

City Analysis & Citizen Services AI: Project Documentation

Project documentation

1.Introduction

Project title : City Analysis & Citizen Services AI Project Documentation

Team ID : NM2025TMID02168

Team Leader : AVINASH R

Team member : KAMALESH S

Team member : KIRUBA SAGAR S

Team member : KAVIRAJ K

2. Project Overview

Purpose:

- * To create an AI-powered assistant that:
 - * Analyzes cities for crime index and accident safety statistics.
 - * Provides overall safety assessments for given cities.
 - * Acts as a government assistant for answering queries on public services, policies, and civic issues.
 - * Builds a user-friendly web application with Gradio for easy interaction.
 - * Demonstrates how Large Language Models (LLMs) like IBM Granite can be applied in governance and safety domains.

Key Features:

1. City Analysis

- * Users input a city name.
- * The AI provides crime index, accident rates, and safety analysis.

2. Citizen Interaction

- * Users input a query related to public services or policies.
- * AI provides detailed, helpful, and accurate responses.

3. Interactive Web Interface (Gradio)

- * Tab-based UI with two sections: City Analysis & Citizen Services.
- * Outputs displayed in large text areas.

Tech Stack:

- * Python (main programming language)
- * Gradio (for interactive web interface)
- * Hugging Face Transformers (IBM Granite model)
- * PyTorch (for GPU/CPU model execution)

3. Architecture

1. Model & Tokenizer Layer

- * Model: ibm-granite/granite-3.2-2b-instruct
- * Library: Hugging Face Transformers (AutoModelForCausalLM, AutoTokenizer)
- Framework: PyTorch

2. Response Generation Layer

- * Function: generate_response()
- * Converts input into prompts → model → readable output.

3. Application Logic Layer

* Functions:

* `city_analysis(city_name)` → Generates analysis for crime index, accidents, safety.

* `citizen_interaction(query)` → Generates policy or service-related responses.

4. User Interface Layer (Gradio)

* Tabs: City Analysis & Citizen Services.

* Inputs: Textbox for city or query.

* Outputs: Large textbox for AI-generated responses.

5. Deployment Layer

* `app.launch(share=True)` → Launches app with a public share link.

■ High-Level Flow (Architecture in Words):

User Input (City Name / Query)

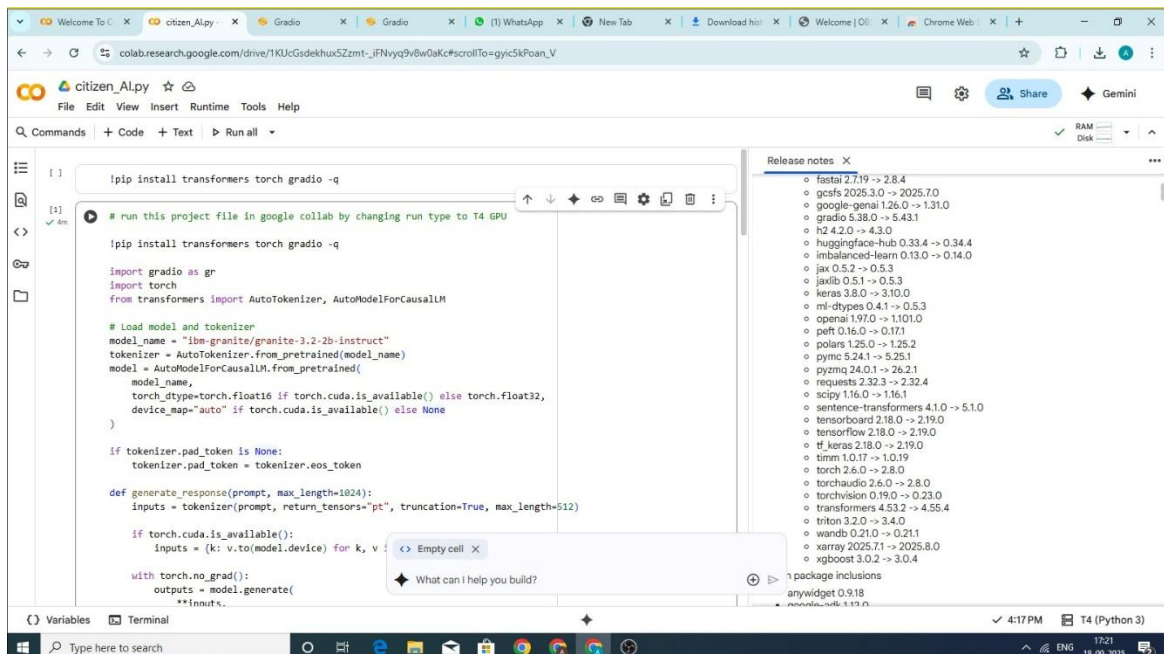
■

▼

Gradio UI (Textbox + Buttons)

SCREEN SHOTS

1.Input:



The screenshot shows a Google Colab notebook titled 'citizen_Alpy'. The code in the first cell is as follows:

```
!pip install transformers torch gradio -q

# run this project file in google collab by changing run type to T4 GPU

!pip install transformers torch gradio -q

import gradio as gr
import torch
from transformers import AutoTokenizer, AutoModelForCausalLM

# Load model and tokenizer
model_name = "ibm-granite/granite-3.2-2b-instruct"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForCausalLM.from_pretrained(
    model_name,
    torch_dtype=torch.float16 if torch.cuda.is_available() else torch.float32,
    device_map="auto" if torch.cuda.is_available() else None
)

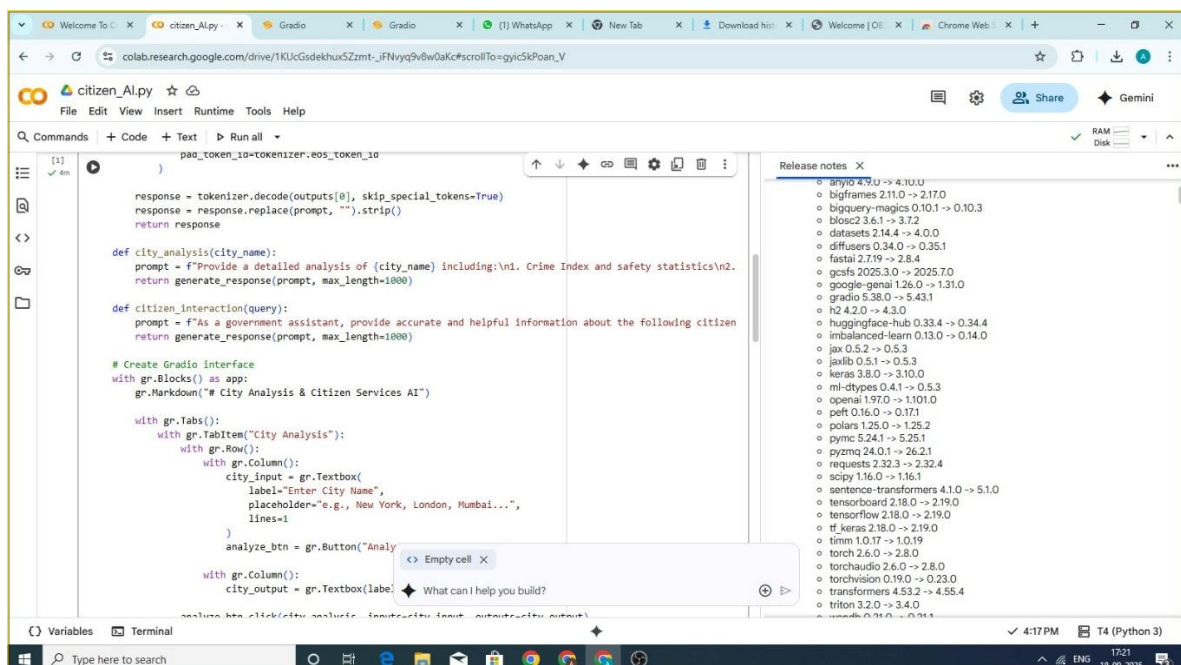
if tokenizer.pad_token is None:
    tokenizer.pad_token = tokenizer.eos_token

def generate_response(prompt, max_length=1024):
    inputs = tokenizer(prompt, return_tensors="pt", truncation=True, max_length=512)

    if torch.cuda.is_available():
        inputs = {k: v.to(model.device) for k, v in inputs.items()}

    with torch.no_grad():
        outputs = model.generate(
            **inputs
        )
```

The right sidebar shows a list of release notes for various packages, including fastai, gcfs, google-gemini, gradio, h2, huggingface-hub, imbalanced-learn, jax, jaxlib, keras, ml-dtypes, openai, peft, polars, pymc, pyzmq, requests, scipy, sentence-transformers, tensorboard, tensorflow, tf.keras, timm, torch, torchaudio, torchvision, transformers, triton, wandb, xarray, and xgboost.



The screenshot shows the same Google Colab notebook, but with the code completed. The code in the second cell is as follows:

```
)

response = tokenizer.decode(outputs[0], skip_special_tokens=True)
response = response.replace(prompt, "").strip()
return response

def city_analysis(city_name):
    prompt = f"Provide a detailed analysis of {city_name} including:\n1. Crime Index and safety statistics\n2.
    return generate_response(prompt, max_length=1000)

def citizen_interaction(query):
    prompt = f"As a government assistant, provide accurate and helpful information about the following citizen
    return generate_response(prompt, max_length=1000)

# Create Gradio Interface
with gr.Blocks() as app:
    gr.Markdown("# City Analysis & Citizen Services AI")

    with gr.Tabs():
        with gr.TabItem("City Analysis"):
            with gr.Row():
                with gr.Column():
                    city_input = gr.Textbox(
                        label="Enter City Name",
                        placeholder="e.g., New York, London, Mumbai...",
                        lines=1
                    )
                    analyze_btn = gr.Button("Analyze")

            with gr.Column():
                city_output = gr.Textbox(label="City Analysis")

        with gr.TabItem("Citizen Interaction"):
            with gr.Row():
                with gr.Column():
                    query_input = gr.Textbox(
                        label="Enter your query",
                        placeholder="e.g., What is the crime index in New York?",
                        lines=1
                    )
                    interact_btn = gr.Button("Interact")

            with gr.Column():
                citizen_output = gr.Textbox(label="Citizen Interaction")

    app.launch()

# Example usage
city_name = "New York"
response = city_analysis(city_name)
print(f"City Analysis: {response}")

query = "What is the crime index in New York?"
response = citizen_interaction(query)
print(f"Citizen Interaction: {response}")
```

The right sidebar shows the same list of release notes as the first screenshot.

OUTPUT:

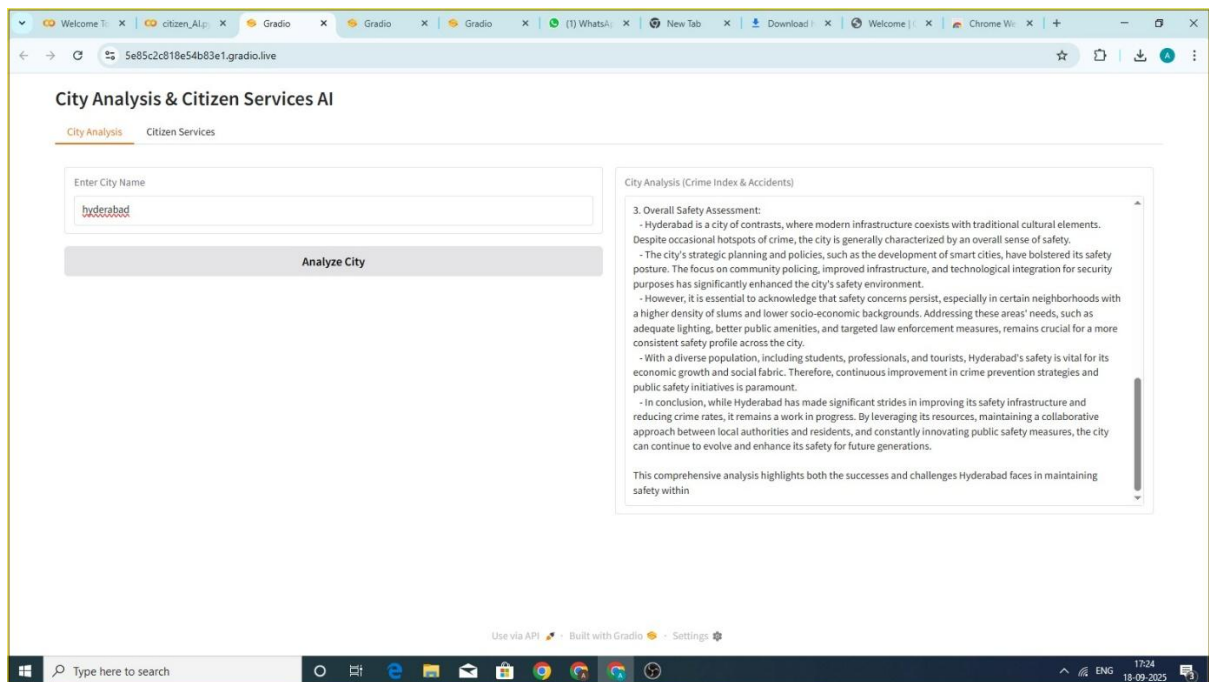
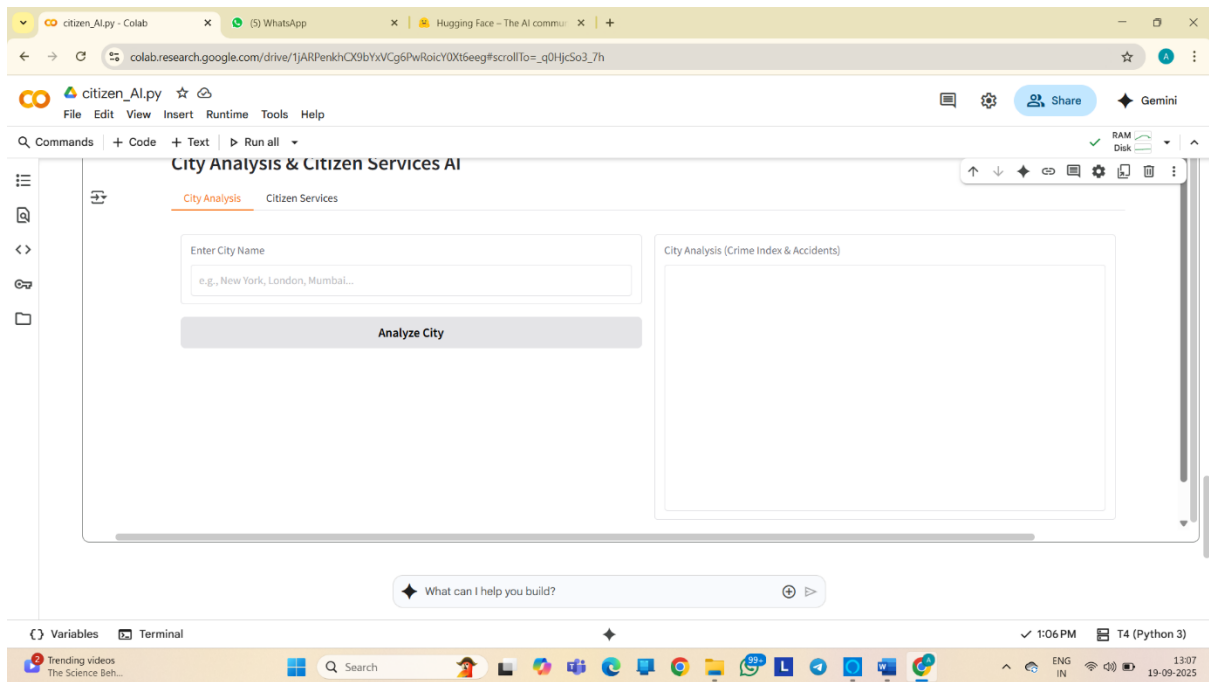
The screenshot displays a Google Colab notebook interface. The notebook is titled "citizen_AI.py" and contains Python code for a web application using Gradio. The code defines a text input for "Your Query", a button "Get Information", and a text output for "Government Response". The application is launched with `app.launch(share=True)`. The execution progress bar shows the following status:

- tokenizer_config.json: 8.88k/? [00:00<00:00, 772kB/s]
- vocab.json: 777k/? [00:00<00:00, 12.3MB/s]
- merges.txt: 442k/? [00:00<00:00, 24.4MB/s]
- tokenizer.json: 3.48M/? [00:00<00:00, 65.6MB/s]
- added_tokens.json: 100% [00:00<00:00, 8.88kB/s]
- special_tokens_map.json: 100% [00:00<00:00, 12.3MB/s]
- config.json: 100% [00:00<00:00, 24.4MB/s]
- torch_dtype: 100% [00:00<00:00, 65.6MB/s]

A warning message is displayed: "The secret 'HF_TOKEN' does not exist in your Colab secrets. To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings). Please note that authentication is recommended but still optional to access public models or datasets." A "What can I help you build?" chat bubble is also visible.

The right sidebar shows a list of installed packages and their versions:

- anyio 4.9.0 -> 4.10.0
- bigframes 2.11.0 -> 2.17.0
- bigquery-magics 0.10.1 -> 0.10.3
- blosc2 3.6.1 -> 3.7.2
- datasets 2.14.4 -> 4.0.0
- diffusers 0.34.0 -> 0.35.1
- fastai 2.7.19 -> 2.8.4
- gcfsfs 2025.3.0 -> 2025.7.0
- google-genai 1.26.0 -> 1.31.0
- gradio 5.38.0 -> 5.43.1
- h2 4.2.0 -> 4.3.0
- huggingface-hub 0.33.4 -> 0.34.4
- imbalanced-learn 0.13.0 -> 0.14.0
- jax 0.5.2 -> 0.5.3
- jaxlib 0.5.1 -> 0.5.3
- keras 3.8.0 -> 3.10.0
- ml-dtypes 0.4.1 -> 0.5.3
- openai 1.97.0 -> 1.101.0
- peft 0.16.0 -> 0.17.1
- polars 1.25.0 -> 1.25.2
- pymc 5.24.1 -> 5.25.1
- pyzmq 24.0.1 -> 26.2.1
- requests 2.32.3 -> 2.32.4
- scipy 1.16.0 -> 1.16.1
- sentence-transformers 4.1.0 -> 5.1.0
- tensorboard 2.18.0 -> 2.19.0
- tensorflow 2.18.0 -> 2.19.0
- tf-keras 2.18.0 -> 2.19.0
- timm 1.0.17 -> 1.0.19
- torch 2.6.0 -> 2.8.0
- torchaudio 2.6.0 -> 2.8.0
- torchvision 0.19.0 -> 0.23.0
- transformers 4.53.2 -> 4.55.4
- triton 3.2.0 -> 3.4.0



Welcome To...citizen_AI.p...GradioGradioGradio(1) WhatsAppNew TabDownloadWelcome | Chrome Wi...+

5e85c2c818e54b83e1.gradio.live

☆🔍⬇️🔒⋮

City Analysis & Citizen Services AI

City AnalysisCitizen Services

Your Query

how to apply for a birth certificate

Get Information

Government Response

- **"In-person"**: You can visit your local vital records office, often located in the county courthouse. Prepare to provide the requested documents and pay the required fees.

- **"Mail"**: Some states accept applications by mail. You'll usually need to include a self-addressed, stamped envelope for the return of your birth certificate.

3. **"Complete the application"**: Fill out the birth certificate request form available through your state's vital records department. You'll need to provide information about yourself and the birth.

4. **"Pay any applicable fees"**: Fees vary by state, but typically range from \$10 to \$40 per certificate. Payment methods usually include cash, money orders, or credit/debit cards.

5. **"Submit your application"**: If you're applying online or by mail, ensure you submit the complete application form along with all required documents and fees.

6. **"Wait for processing"**: Processing times can range from a few weeks to several months, depending on the state's workload and the method of application.

7. **"Receive your certificate"**: Once your application is approved, you'll receive your birth certificate by mail.

For the most accurate and up-to-date information, visit your state's official vital records website or contact the local vital records office in your county. Remember that processes and specific requirements can differ slightly.

Use via API - Built with Gradio - Settings

Type here to search

ENG17:2418-09-2025